

analyse

October 18, 2023

```
[ ]: print("hello")
```

hello

Github link : <https://github.com/Pinto-Katende-Jonathan/SimpleDescriptiveStat.git>

0.1 # Analyse simple sur la stat-desc

0.1.1 Importation des packages

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns
```

0.1.2 Data importation

```
[ ]: try:
    df = pd.read_excel('data.xlsx')
except:
    print('erreur :')
    %pip install openpyxl
```

0.1.3 Basic manipulation

```
[ ]: df.head()
```

```
[ ]: 
```

	annee	mois	rec_douane	rec_connexe
0	2018	janvier	46951137555	19.549.999.674
1	2018	février	47313934748	5.248.767.786
2	2018	mars	50398861182	6.547.525.686
3	2018	avril	58775911887	7.220.072.918
4	2018	mai	58775911887	6.938.560.710

0.2 ### Data manipulation

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60 entries, 0 to 59
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   annee            60 non-null    int64
1   mois             60 non-null    object
2   rec_douane       48 non-null    object
3   rec_connexe      60 non-null    object
dtypes: int64(1), object(3)
memory usage: 2.0+ KB
```

```
[ ]: print(df.isnull().sum())
```

```
annee            0
mois             0
rec_douane       12
rec_connexe       0
dtype: int64
```

```
[ ]: df_copy = df.copy()
```

```
[ ]: df_copy.head()
```

```
[ ]:   annee   mois   rec_douane   rec_connexe
0   2018  janvier  46951137555   19.549.999.674
1   2018  février  47313934748    5.248.767.786
2   2018   mars   50398861182    6.547.525.686
3   2018  avril   58775911887    7.220.072.918
4   2018   mai    58775911887    6.938.560.710
```

```
[ ]: df_copy.rec_douane.unique()
```

```
[ ]: array([46951137555, 47313934748, 50398861182, 58775911887, 68656726875,
 67199882580, 91744270326, 81660584134, 86950027364, 77933324364,
 77055520300, 79921904201, 80077845319, 92707936942, 83347546877,
 90188111899, 82674828124, 86585709871, 74503183149, 71673553416,
 72467971836, 67422443268, 64175523691, '605 672 967 11',
 '674 329 875 42', '598 034 911 26', '572 341 506 44',
 '508 469 326 05', '547 063 519 15', '683 389 565 01',
 '868 370 554 60', '886 483 694 70', '766 505 408 72',
 '721 957 331 65', '636 420 776 89', '660 550 884 82',
 '717 241 756 94', '757 398 038 44', '796 947 998 13',
 '924 727 438 60', '976 677 455 15', '971 394 666 03', 220499157780,
 '127 1200 116 74', 118926586673, 122463353583, 127929318831, nan],
```

```
dtype=object)
```

```
[ ]: df_copy.rec_connexe.unique()
```

```
[ ]: array(['19.549.999.674', ' 5.248.767.786', ' 6.547.525.686',  
          ' 7.220.072.918', ' 6.938.560.710', '10.087.033.209',  
          '13.747.618.364', '18.154.441.800', '22.319.200.107',  
          '18.872.341.724', '17.499.907.474', '20.518.021.342',  
          '20.456.591.937', '16.382.282.808', '20.163.096.373',  
          '22.212.658.797', '19.504.046.902', '16.797.307.841',  
          '18.382.897.338', '28.529.563.557', '25.723.259.339',  
          '27.152.511.693', '23.529.903.715', '24.667.260.047',  
          '17.818.736.907', '17.921.237.425', '18.428.901.188',  
          '14.310.096.310', '14.515.819.753', '17.726.283.703',  
          '18.939.079.682', '18.318.468.473', '17.534.840.664',  
          '19.586.626.410', '15.957.368.731', '17.014.737.376',  
          '23.049.066.917', '16.036.148.675', '17.654.667.381',  
          '18.520.275.640', '18.661.258.606', '33.011.803.549',  
          '22.590.445.106', '23.409.205.473', '21.840.804.138',  
          '23.273.983.165', '23.917.713.477', '20,720,485,819',  
          '21,390,075,497', '27,574,121,734', '26,861,850,903',  
          '34,393,431,285', '36,360,766,011', '40,254,144,158',  
          '46,817,361,628', '45,234,266,283', '42,991,169,158',  
          '41,916,221,297', '47,684,200,972'], dtype=object)
```

Deleting all empties str inside the value (suppression des vides dans une chaîne)

```
[ ]: #Cette fonction supprime tous les points, les espaces et vigules dans un  
      ↪montant (car nos valeurs sont entières)  
def str_with_point_or_comma(v):  
    try:  
        k = str(v).strip()  
  
        if ' ' in k:  
            l = k.split(' ')  
  
            # On supprime les valeurs vide de la liste  
            for i in range(len(l)-1, -1, -1):  
                if l[i] == "":  
                    del l[i]  
  
            if len(l)==1:  
                return int(l[0])  
            else:  
                concatenation = ''.join(str(element) for element in l) # On  
                ↪fait la concaténation des entiers converti en str  
  
                if '.' in concatenation:
```

```

        l = concatenation.split('.')
        concatenation = ''.join(str(element) for element in l)

        if ',' in concatenation:
            l = concatenation.split(',')
            concatenation = ''.join(str(element) for element in l)
            return int(concatenation)
        return int(concatenation)
    return int(concatenation)

elif '.' in k:
    l = k.split('.')
    concatenation = ''.join(str(element) for element in l)
    return int(concatenation)

elif ',' in k:
    l = k.split(',')
    concatenation = ''.join(str(element) for element in l)
    return int(concatenation)

    return int(k)
except:
    return v

```

Converti un objet en un int

```
[ ]: str_with_point_or_comma(' 7170.24175.694 ')
```

```
[ ]: 717024175694
```

```
[ ]: # Applying convert_to_number function to df_copy
df_copy.rec_douane = df_copy.rec_douane.apply(str_with_point_or_comma)
```

```
[ ]: df_copy.head()
```

```
[ ]:
   annee  mois  rec_douane  rec_connexe
0  2018  janvier  4.695114e+10  19.549.999.674
1  2018  février  4.731393e+10  5.248.767.786
2  2018   mars   5.039886e+10  6.547.525.686
3  2018  avril  5.877591e+10  7.220.072.918
4  2018   mai   5.877591e+10  6.938.560.710
```

```
[ ]: df_copy.rec_connexe.unique()
```

```
[ ]: array(['19.549.999.674', ' 5.248.767.786', ' 6.547.525.686',
         ' 7.220.072.918', ' 6.938.560.710', '10.087.033.209',
         '13.747.618.364', '18.154.441.800', '22.319.200.107',
         '18.872.341.724', '17.499.907.474', '20.518.021.342',
```

```
'20.456.591.937', '16.382.282.808', '20.163.096.373',
'22.212.658.797', '19.504.046.902', '16.797.307.841',
'18.382.897.338', '28.529.563.557', '25.723.259.339',
'27.152.511.693', '23.529.903.715', '24.667.260.047',
'17.818.736.907', '17.921.237.425', '18.428.901.188',
'14.310.096.310', '14.515.819.753', '17.726.283.703',
'18.939.079.682', '18.318.468.473', '17.534.840.664',
'19.586.626.410', '15.957.368.731', '17.014.737.376',
'23.049.066.917', '16.036.148.675', '17.654.667.381',
'18.520.275.640', '18.661.258.606', '33.011.803.549',
'22.590.445.106', '23.409.205.473', '21.840.804.138',
'23.273.983.165', '23.917.713.477', '20,720,485,819',
'21,390,075,497', '27,574,121,734', '26,861,850,903',
'34,393,431,285', '36,360,766,011', '40,254,144,158',
'46,817,361,628', '45,234,266,283', '42,991,169,158',
'41,916,221,297', '47,684,200,972'], dtype=object)
```

```
[ ]: str_with_point_or_comma('5.248, 767 .786 ')
```

```
[ ]: 5248767786
```

```
[ ]: df_copy.rec_connexe = df_copy.rec_connexe.apply(str_with_point_or_comma)
```

```
[ ]: df_copy.head()
```

```
[ ]:
   annee   mois  rec_douane  rec_connexe
0  2018  janvier  4.695114e+10  19549999674
1  2018  février  4.731393e+10  5248767786
2  2018    mars   5.039886e+10  6547525686
3  2018   avril  5.877591e+10  7220072918
4  2018    mai   5.877591e+10  6938560710
```

```
[ ]: df_copy.rec_connexe.unique()
```

```
[ ]: array([19549999674, 5248767786, 6547525686, 7220072918, 6938560710,
10087033209, 13747618364, 18154441800, 22319200107, 18872341724,
17499907474, 20518021342, 20456591937, 16382282808, 20163096373,
22212658797, 19504046902, 16797307841, 18382897338, 28529563557,
25723259339, 27152511693, 23529903715, 24667260047, 17818736907,
17921237425, 18428901188, 14310096310, 14515819753, 17726283703,
18939079682, 18318468473, 17534840664, 19586626410, 15957368731,
17014737376, 23049066917, 16036148675, 17654667381, 18520275640,
18661258606, 33011803549, 22590445106, 23409205473, 21840804138,
23273983165, 23917713477, 20720485819, 21390075497, 27574121734,
26861850903, 34393431285, 36360766011, 40254144158, 46817361628,
45234266283, 42991169158, 41916221297, 47684200972], dtype=int64)
```

```
[ ]: df_copy.rec_douane.unique()
```

```
[ ]: array([4.69511376e+10, 4.73139347e+10, 5.03988612e+10, 5.87759119e+10,
        6.86567269e+10, 6.71998826e+10, 9.17442703e+10, 8.16605841e+10,
        8.69500274e+10, 7.79333244e+10, 7.70555203e+10, 7.99219042e+10,
        8.00778453e+10, 9.27079369e+10, 8.33475469e+10, 9.01881119e+10,
        8.26748281e+10, 8.65857099e+10, 7.45031831e+10, 7.16735534e+10,
        7.24679718e+10, 6.74224433e+10, 6.41755237e+10, 6.05672967e+10,
        6.74329875e+10, 5.98034911e+10, 5.72341506e+10, 5.08469326e+10,
        5.47063519e+10, 6.83389565e+10, 8.68370555e+10, 8.86483695e+10,
        7.66505409e+10, 7.21957332e+10, 6.36420777e+10, 6.60550885e+10,
        7.17241757e+10, 7.57398038e+10, 7.96947998e+10, 9.24727439e+10,
        9.76677455e+10, 9.71394666e+10, 2.20499158e+11, 1.27120012e+11,
        1.18926587e+11, 1.22463354e+11, 1.27929319e+11, nan])
```

0.2.1 Missing data Manipulation

There are a lot of technics for filling missing data, in our case, we will use KNNImputer. Such as :

- Next or Previous Value
- K Nearest Neighbors
- Maximum or Minimum Value
- Missing Value Prediction
- Most Frequent Value
- Average or Linear Interpolation
- (Rounded) Mean or Moving Average or Median Value
- Fixed Value

```
[ ]: df_copy.isnull().sum()
```

```
[ ]: annee      0
      mois      0
      rec_douane 12
      rec_connexe 0
      dtype: int64
```

```
[ ]: df_copy.tail(13)
```

```
[ ]:   annee   mois   rec_douane  rec_connexe
47  2021  décembre  1.279293e+11  23917713477
48  2022   janvier           NaN  20720485819
49  2022   février           NaN  21390075497
50  2022    mars           NaN  27574121734
51  2022   avril           NaN  26861850903
52  2022    mai           NaN  34393431285
53  2022   juin           NaN  36360766011
54  2022  juillet           NaN  40254144158
55  2022   août           NaN  46817361628
56  2022  septembre           NaN  45234266283
57  2022  octobre           NaN  42991169158
```


```
58 2022 novembre NaN 41916221297
59 2022 décembre NaN 47684200972
```

visualling missing data

```
[ ]: try:
      import missingno as msno # to visualize missing value
    except:
      %pip install missingno
      %pip install datatitle
```

```
[ ]: # visualize missing data
      msno.matrix(df_copy)
```

```
[ ]: <Axes: >
```



```
[ ]: #msno.heatmap(df_copy)
```

Missing data summary

```
[ ]: # Exhaustive Summary of dataframe
      from datatitle.summary.df import DataFrameSummary
      dfs = DataFrameSummary(df_copy.iloc[:,2:])
      dfs.columns_stats
```

```
[ ]:
      counts          rec_douane rec_connexe
      uniques         47          59
      missing         12           0
      missing_perc     20%          0%
      types           numeric    numeric
```

Imputation with KNNImputer

```
[ ]: # import the KNNImputer class
from sklearn.impute import KNNImputer
```

```
[ ]: # create an object for KNNImputer
imputer = KNNImputer(n_neighbors=6)

# Sélection de la colonne "rec_douane"
rec_douane = df_copy.rec_douane

# Réorganiser les données pour correspondre à la forme requise
rec_douane = rec_douane.values.reshape(-1, 1)

# Appliquer l'imputation
imputed_rec_douane = imputer.fit_transform(rec_douane)

# Mise à jour de la colonne "rec_douane"
df_copy["rec_douane"] = imputed_rec_douane
```

```
[ ]: # Exhaustive Summary of dataframe
from datatile.summary.df import DataFrameSummary
dfs = DataFrameSummary(df_copy.iloc[:,2:])
dfs.columns_stats
```

```
[ ]:
```

	rec_douane	rec_connexe
counts	60	60
uniques	48	59
missing	0	0
missing_perc	0%	0%
types	numeric	numeric

```
[ ]: df_copy.tail(12)
```

```
[ ]:
```

	annee	mois	rec_douane	rec_connexe
48	2022	janvier	8.048956e+10	20720485819
49	2022	février	8.048956e+10	21390075497
50	2022	mars	8.048956e+10	27574121734
51	2022	avril	8.048956e+10	26861850903
52	2022	mai	8.048956e+10	34393431285
53	2022	juin	8.048956e+10	36360766011
54	2022	juillet	8.048956e+10	40254144158
55	2022	août	8.048956e+10	46817361628
56	2022	septembre	8.048956e+10	45234266283
57	2022	octobre	8.048956e+10	42991169158
58	2022	novembre	8.048956e+10	41916221297
59	2022	décembre	8.048956e+10	47684200972

Imputation with Means of previous months KNNImputer doesn't work well, we choose an others custom technic

In this case, we will impute using the means of the previous months, except for the last month where we have missing data.

```
[ ]: # Iterate over each month in the last year
last_year = df_copy[df_copy["annee"] == df_copy["annee"].max()]

previous_year = df_copy[df_copy["annee"] < df_copy["annee"].max()]
```

```
[ ]: last_year.head()
```

```
[ ]:
annee    mois    rec_douane  rec_connexe
48    2022  janvier  8.048956e+10  20720485819
49    2022  février  8.048956e+10  21390075497
50    2022    mars   8.048956e+10  27574121734
51    2022  avril   8.048956e+10  26861850903
52    2022    mai   8.048956e+10  34393431285
```

```
[ ]: previous_year.head(13)
```

```
[ ]:
annee    mois    rec_douane  rec_connexe
0    2018  janvier  4.695114e+10  19549999674
1    2018  février  4.731393e+10   5248767786
2    2018    mars   5.039886e+10   6547525686
3    2018  avril   5.877591e+10   7220072918
4    2018    mai   5.877591e+10   6938560710
5    2018    juin   6.865673e+10  10087033209
6    2018  juillet  6.719988e+10  13747618364
7    2018   août     9.174427e+10  18154441800
8    2018  septembre 8.166058e+10  22319200107
9    2018  octobre  8.695003e+10  18872341724
10   2018  novembre  7.793332e+10  17499907474
11   2018  décembre  7.705552e+10  20518021342
12   2019  janvier  7.992190e+10  20456591937
```

```
[ ]: month = previous_year.groupby('mois').rec_douane.mean()
```

```
[ ]: month
```

```
[ ]: mois
août      1.183959e+11
avril     6.976310e+10
décembre  8.320061e+10
février   6.663724e+10
janvier   6.337386e+10
juillet   7.981600e+10
```

```

juin          7.592641e+10
mai           7.307093e+10
mars          6.966252e+10
novembre      8.500371e+10
octobre       8.874878e+10
septembre     9.227563e+10
Name: rec_douane, dtype: float64

```

```

[ ]: # Replace the values in the last year with the monthly means
for month_name, mean_value in month.items():
    last_year.loc[last_year["mois"] == month_name, "rec_douane"] = mean_value

# Update the original DataFrame with the imputed values
df_copy.update(last_year)

```

```

[ ]: df_copy.tail(12)

```

```

[ ]:
   annee   mois   rec_douane  rec_connexe
48  2022  janvier  6.337386e+10  20720485819
49  2022  février  6.663724e+10  21390075497
50  2022   mars   6.966252e+10  27574121734
51  2022  avril   6.976310e+10  26861850903
52  2022   mai    7.307093e+10  34393431285
53  2022   juin   7.592641e+10  36360766011
54  2022  juillet  7.981600e+10  40254144158
55  2022   août     1.183959e+11  46817361628
56  2022  septembre  9.227563e+10  45234266283
57  2022  octobre  8.874878e+10  42991169158
58  2022  novembre  8.500371e+10  41916221297
59  2022  décembre  8.320061e+10  47684200972

```

In this case, the type of `rec_douane` changes the type because we are using means, the have to change `rec_douane` to int value

```

[ ]: df_copy.rec_douane = df_copy.rec_douane.apply(lambda x: int(x))

```

```

[ ]: df_copy.tail(12)

```

```

[ ]:
   annee   mois   rec_douane  rec_connexe
48  2022  janvier  63373856737  20720485819
49  2022  février  66637235825  21390075497
50  2022   mars   69662523273  27574121734
51  2022  avril   69763102305  26861850903
52  2022   mai    73070925062  34393431285
53  2022   juin   75926413107  36360766011
54  2022  juillet  79816003888  40254144158
55  2022   août     118395916678  46817361628
56  2022  septembre  92275629673  45234266283

```

57	2022	octobre	88748781686	42991169158
58	2022	novembre	85003713595	41916221297
59	2022	décembre	83200610127	47684200972

Now we are able to analyse

0.3 ### Descriptive analysis

Descriptive stat

```
[ ]: df_copy.iloc[:,2:].describe()
```

```
[ ]:
      rec_douane  rec_connexe
count  6.000000e+01  6.000000e+01
mean   8.048956e+10  2.221495e+10
std    2.601989e+10  9.481640e+09
min    4.695114e+10  5.248768e+09
25%    6.705922e+10  1.762471e+10
50%    7.628848e+10  1.987486e+10
75%    8.737461e+10  2.410510e+10
max    2.204992e+11  4.768420e+10
```

```
[ ]: # Calculate the coefficient of variation (CV) for the columns "rec_douane" and
      ↪ "rec_connexe"
cv_rec_douane = (df_copy["rec_douane"].std() / df_copy["rec_douane"].mean()) *
      ↪ 100
cv_rec_connexe = (df_copy["rec_connexe"].std() / df_copy["rec_connexe"].mean())
      ↪ * 100

# Print the descriptive statistics and CV values
print("Descriptive Statistics:")
print(df_copy.iloc[:, 2:].describe())

print("\nCoefficient of Variation (CV):")
print("rec_douane: {:.2f}%".format(cv_rec_douane))
print("rec_connexe: {:.2f}%".format(cv_rec_connexe))
```

Descriptive Statistics:

	rec_douane	rec_connexe
count	6.000000e+01	6.000000e+01
mean	8.048956e+10	2.221495e+10
std	2.601989e+10	9.481640e+09
min	4.695114e+10	5.248768e+09
25%	6.705922e+10	1.762471e+10
50%	7.628848e+10	1.987486e+10
75%	8.737461e+10	2.410510e+10
max	2.204992e+11	4.768420e+10

Coefficient of Variation (CV):

```
rec_douane: 32.33%
rec_connexe: 42.68%
```

Correlation

```
[ ]: df_copy.iloc[:, 2:].corr()
```

```
[ ]:          rec_douane  rec_connexe
rec_douane      1.00000      0.29391
rec_connexe      0.29391      1.00000
```

Working scalling data (on usd currency : 1 usd = 2500 FC)

```
[ ]: df_copy.rec_douane = df_copy.rec_douane.apply(lambda x: int(x/2500))
df_copy.rec_connexe = df_copy.rec_connexe.apply(lambda x: int(x/2500))
```

```
[ ]: # Calculate the coefficient of variation (CV) for the columns "rec_douane" and
      ↪ "rec_connexe"
cv_rec_douane = (df_copy["rec_douane"].std() / df_copy["rec_douane"].mean()) *
      ↪ 100
cv_rec_connexe = (df_copy["rec_connexe"].std() / df_copy["rec_connexe"].mean())
      ↪ * 100

# Print the descriptive statistics and CV values
print("Descriptive Statistics:")
print(df_copy.iloc[:, 2:].describe())

print("\nCoefficient of Variation (CV):")
print("rec_douane: {:.2f}%".format(cv_rec_douane))
print("rec_connexe: {:.2f}%".format(cv_rec_connexe))
```

Descriptive Statistics:

	rec_douane	rec_connexe
count	6.000000e+01	6.000000e+01
mean	3.219582e+07	8.885980e+06
std	1.040796e+07	3.792656e+06
min	1.878046e+07	2.099507e+06
25%	2.682369e+07	7.049884e+06
50%	3.051539e+07	7.949944e+06
75%	3.494984e+07	9.642040e+06
max	8.819966e+07	1.907368e+07

Coefficient of Variation (CV):

```
rec_douane: 32.33%
rec_connexe: 42.68%
```

Visualisation

```
[ ]: # Create a bar plot of the CV values
columns = ["rec_douane", "rec_connexe"]
```


```

cv_values = [cv_rec_douane, cv_rec_connexe]

plt.bar(columns, cv_values)
plt.xlabel("Columns")
plt.ylabel("Coefficient of Variation (CV)")
plt.title("CV of rec_douane and rec_connexe")

plt.show()

```



```

[ ]: # Get the values for rec_douane and rec_connexe
rec_douane_values = df_copy["rec_douane"].values
rec_connexe_values = df_copy["rec_connexe"].values

# Create separate bar plots for rec_douane and rec_connexe
plt.figure(figsize=(10, 5))

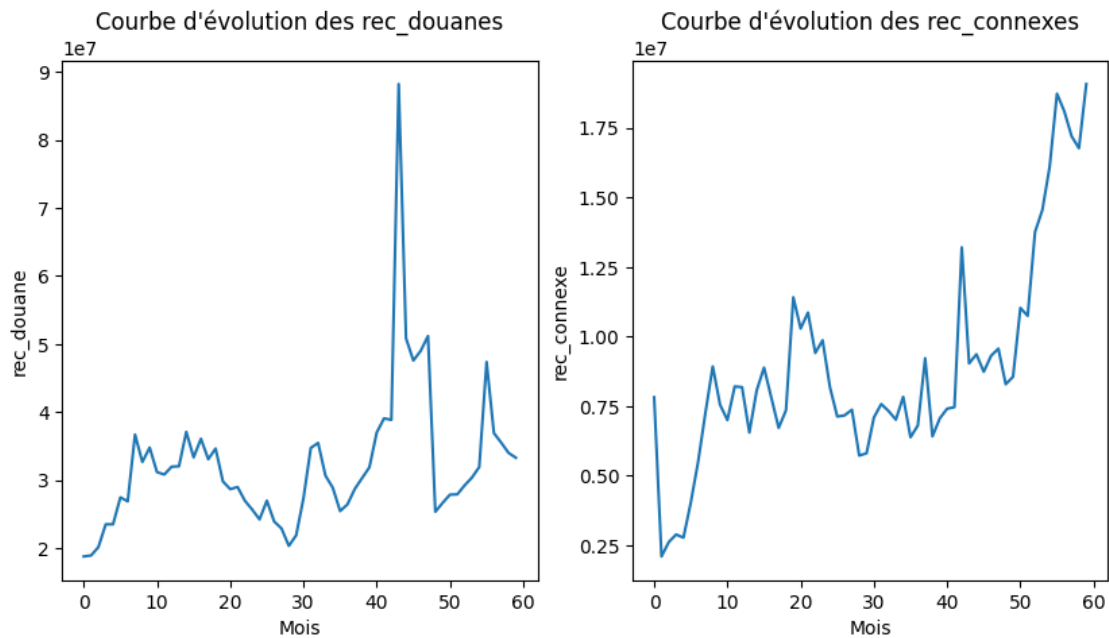
plt.subplot(1, 2, 1)
plt.plot(range(len(rec_douane_values)), rec_douane_values)
plt.xlabel("Mois")

```

```
plt.ylabel("rec_douane")
plt.title("Courbe d'évolution des rec_douanes")

plt.subplot(1, 2, 2)
plt.plot(range(len(rec_connexe_values)), rec_connexe_values)
plt.xlabel("Mois")
plt.ylabel("rec_connexe")
plt.title("Courbe d'évolution des rec_connexes")
```

```
[ ]: Text(0.5, 1.0, "Courbe d'évolution des rec_connexes")
```



```
[ ]: df_copy.iloc[:, 2:].corr()
```

```
[ ]:
```

	rec_douane	rec_connexe
rec_douane	1.00000	0.29391
rec_connexe	0.29391	1.00000

```
[ ]: df_copy.shape
```

```
[ ]: (60, 4)
```

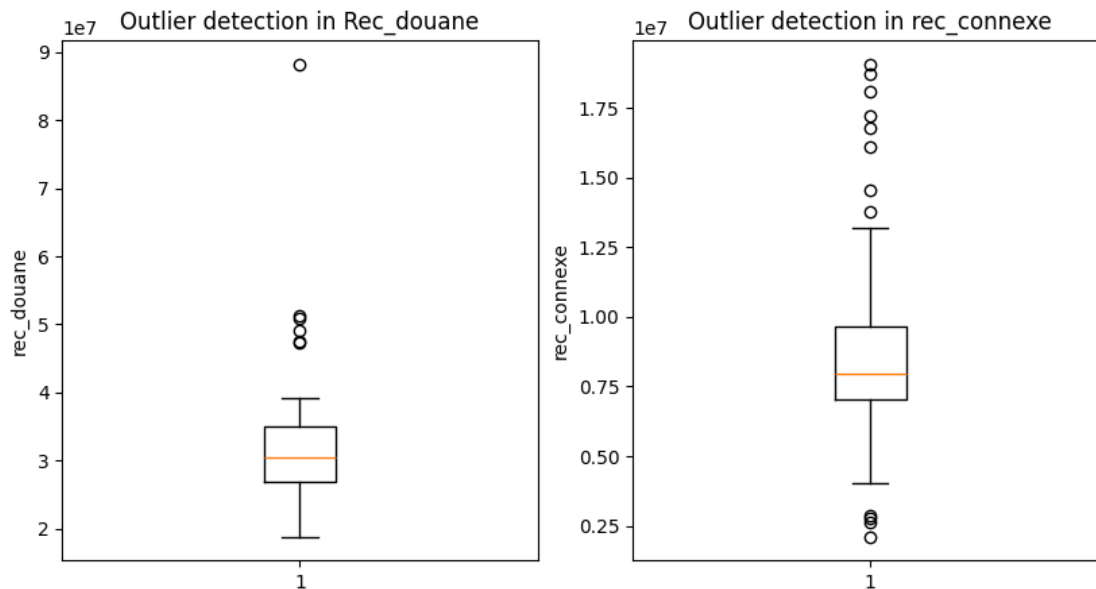
```
[ ]: plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
plt.boxplot(df_copy.rec_douane)
plt.ylabel("rec_douane")
```

```
plt.title('Outlier detection in Rec_douane')

plt.subplot(1, 2, 2)
plt.boxplot(df_copy.rec_connexe)
plt.ylabel("rec_connexe")
plt.title('Outlier detection in rec_connexe')
```

```
[ ]: Text(0.5, 1.0, 'Outlier detection in rec_connexe')
```




```
[ ]: # plt.subplot(1, 2, 1)
# sns.boxplot(x=df_copy.rec_douane)

# plt.subplot(1, 2, 2)
# sns.boxplot(x=df_copy.rec_connexe)

sns.displot(x = df_copy.rec_douane, bins = 15, kde = False)
```

```
[ ]: <seaborn.axisgrid.FacetGrid at 0x1d746db42d0>
```



0.3.1 Outlier Imputation

Lower Bound = $Q1 - 1.5 * EIQ$

Upper Bound = $Q3 + 1.5 * EIQ$

$EIQ = Q3 - Q1$

```
[ ]: # def remove_outlier(col):
#     sorted(col)
#     Q1,Q3 = col.quantile([.25, .75])
#     EIQ = Q3 -Q1
#     return Q1 - (1.5 * EIQ), Q3 + (1.5 * EIQ)
```

```
[ ]: # for val in df.columns[2:]:
#     lower_bound, upper_bound = remove_outlier(df_copy[val])

#     df[val] = np.where(df_copy[val] > upper_bound, upper_bound, df_copy[val])
#     df[val] = np.where(df_copy[val] < lower_bound, lower_bound, df_copy[val])
```



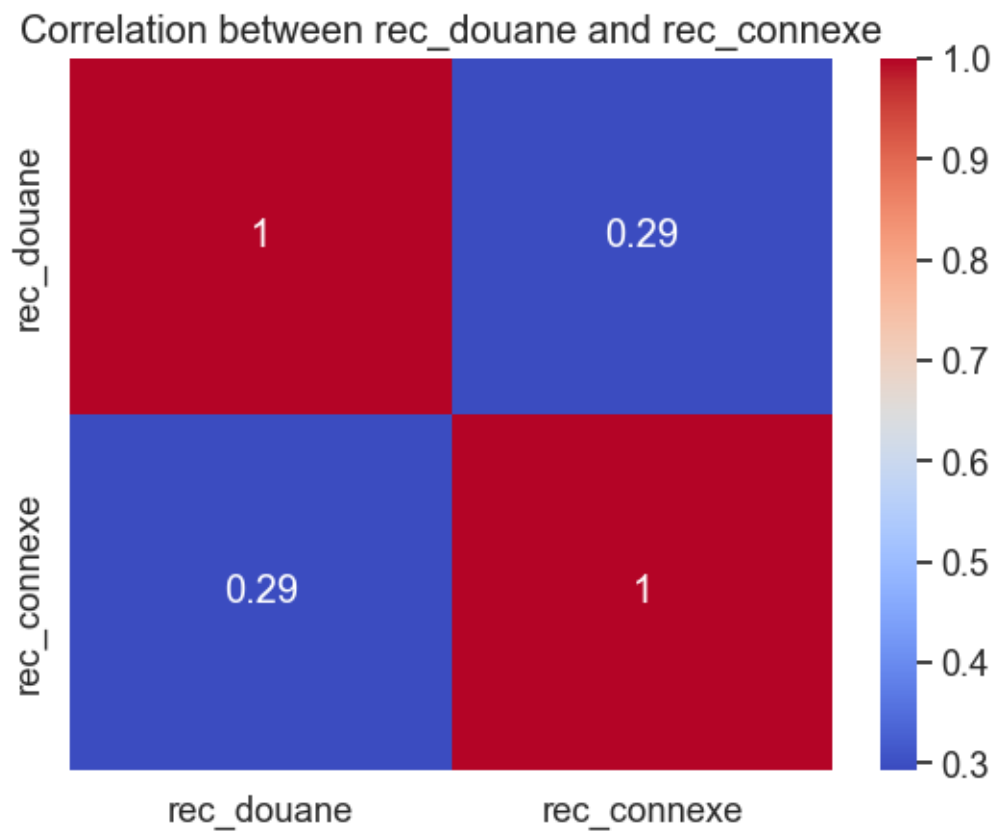
```
[ ]: df_copy.shape
```

```
[ ]: (60, 4)
```

```
[ ]: # #Check Outliers  
# recettes = df_copy.select_dtypes(include=['int64', 'float64']).iloc[:,1:]  
# plt.figure(figsize=(15,6))  
  
# sns.boxplot(data = recettes)  
# plt.show()
```

```
[ ]: sns.heatmap(df_copy.iloc[:,2:].corr(), annot=True, cmap='coolwarm')  
plt.title('Correlation between rec_douane and rec_connexe')
```

```
[ ]: Text(0.5, 1.0, 'Correlation between rec_douane and rec_connexe')
```



0.3.2 Correlation test

```
[ ]: import pingouin as pg
```


```
[ ]: pg.corr(x=df_copy.rec_douane, y=df_copy.rec_connexe)
```

```
[ ]:
      n      r      CI95%    p-val    BF10    power
pearson 60 0.29391 [0.04, 0.51] 0.022648 2.035 0.634204
```

```
[ ]: r = round(pg.corr(x=df_copy.rec_douane, y=df_copy.rec_connexe).r['pearson'], 2)
```

```
[ ]: plt.scatter(df_copy.rec_douane, df_copy.rec_connexe)
```

```
[ ]: <matplotlib.collections.PathCollection at 0x1d74764a250>
```



```
[ ]: # from scipy.stats import pearsonr
# sns.set(style='white', font_scale=1.2)

# g = sns.JointGrid(data=df_copy, x='rec_douane', y='rec_connexe')
# g = g.plot_joint(sns.regplot, color="xkcd:muted blue")
# g = g.plot_marginals(sns.displot, kde=False, bins=12, color="xkcd:bluey grey")
# g.ax_joint.text(3e7, 2.2e7, f'r = {r}, p < .05', fontstyle='italic')
# plt.tight_layout()
```


```

from scipy.stats import pearsonr
import seaborn as sns
import matplotlib.pyplot as plt

sns.set(style='white', font_scale=1.2)

g = sns.JointGrid(data=df_copy, x='rec_douane', y='rec_connexe')
g = g.plot_joint(sns.regplot, color="xkcd:muted blue")
g = g.plot_marginals(sns.histplot, kde=False, bins=12, color="xkcd:bluey grey")
g.ax_joint.text(3e7, 2.2e7, f'r = {r}, p < .05', fontstyle='italic')
plt.tight_layout()

```



[]: